

```
#Amruta Sool
#BBC019128
```

```
!/usr/local/cuda/bin/nvcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2022 NVIDIA Corporation
Built on Wed_Sep_21_10:33:58_PDT_2022
Cuda compilation tools, release 11.8, V11.8.89
Build cuda_11.8.r11.8/compiler.31833905_0
```

```
!nvcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2022 NVIDIA Corporation
Built on Wed_Sep_21_10:33:58_PDT_2022
Cuda compilation tools, release 11.8, V11.8.89
Build cuda_11.8.r11.8/compiler.31833905_0
```

```
!pip install git+https://github.com/andreinechaev/nvcc4jupyter.git
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting git+https://github.com/andreinechaev/nvcc4jupyter.git
  Cloning https://github.com/andreinechaev/nvcc4jupyter.git to /tmp/pip-req-build-bmxxsg23
  Running command git clone --filter=blob:none --quiet https://github.com/andreinechaev/nvcc4jupyter.git /tmp/pip-req-build-bmxxsg23
  Resolved https://github.com/andreinechaev/nvcc4jupyter.git to commit aac710a35f52bb78ab34d2e52517237941399eff
  Preparing metadata (setup.py) ... done
Building wheels for collected packages: NVCCPlugin
  Building wheel for NVCCPlugin (setup.py) ... done
  Created wheel for NVCCPlugin: filename=NVCCPlugin-0.0.2-py3-none-any.whl size=4287 sha256=14ac02382d330fce4ea106136d640a8e5697cc91a62d5a9cfeaf84
  Stored in directory: /tmp/pip-ephem-wheel-cache-_gkvm18c/wheels/a8/b9/18/23f8ef71ceb0f63297dd1903aedd067e6243a68ea756d6f6ea
Successfully built NVCCPlugin
Installing collected packages: NVCCPlugin
Successfully installed NVCCPlugin-0.0.2
```

```
# Commented out IPython magic to ensure Python compatibility.
%load_ext nvcc_plugin
```

```
created output directory at /content/src
Out bin /content/result.out
```

```
# Commented out IPython magic to ensure Python compatibility.
%%cu
```

```
/*[m,n][n,j]= [m,j]
# first take transpose of vector anfd
# then mulriply and add row wisw for each col.
# */
# //parallelism of multiplication only
# //tested ok on 05/12/2018

#include <cuda.h>
#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#define m 10
__global__ void mul_r(int *a, int *b, int *c){
#
int tid = threadIdx.x;
    if (tid < m){
        c[tid]= a[tid] * b[tid];
    }

}

int main(){
    int n, c, d, fst[10][10], snd[10][10], t_snd[10][10];
    int row,col,sum_c, a[10], b[10], ans[10];

/*
    printf("Enter the number of rows and columns of matrix\n");
    scanf("%d%d", &m, &n);
*/
n=m; //square matrix only
    //printf("Enter the elements of first matrix\n");
    //for true random value of vector
    //srand(time(0));
    for (c = 0; c < m; c++)
    {

        for (d = 0; d < n; d++)
        {
```

```

//scanf("%d", &first[c][d]);

fst[c][d]=rand()%10+1;
}
}
printf("display the elements of first matrix\n");
for (c = 0; c < m; c++) {
    for (d = 0 ; d < n; d++) {

        printf("%d\t", fst[c][d]);
        }
        printf("\n");
    }

// take next matrix
//for true random value of vector
//srand(time(0));
for (c = 0; c < m; c++)
{
    for (d = 0; d < n; d++)
    {
        //scanf("%d", &first[c][d]);

        snd[c][d]=rand()%10+1;
    }
}
printf("display the elements of second matrix\n");
for (c = 0; c < m; c++) {
    for (d = 0 ; d < n; d++) {

        printf("%d\t", snd[c][d]);
        }
        printf("\n");
    }
// transpose of second matrix
for(c=0; c<m; c++)
    for(d=0; d<n; d++)
    {
        t_snd[d][c] = snd[c][d];
    }

// Displaying the transpose of matrix a
printf("\nTranspose of second Matrix:\n");
for (c = 0; c < n; c++) {
    for (d = 0 ; d < m; d++) {

        printf("%d\t", t_snd[c][d]);

        }
        printf("\n");
    }

// now multiply on cuda
int *dev_a, *dev_b,*dev_ans;
cudaError_t err=cudaSuccess;
// allocate memory on GPU
err=cudaMalloc((void**)&dev_a,m * sizeof(int));

if (err !=cudaSuccess)
{
    printf("failed to allocate on device \n");
    printf("error is:\n",cudaGetErrorString(err));
    exit(EXIT_FAILURE);
}
//printf("first ok\n");
cudaMalloc((void**)&dev_b,m * sizeof(int));
cudaMalloc((void**)&dev_ans,m * sizeof(int));
//printf("first finished ok\n");
row=0;
col=0;
cudaEvent_t start, end;
cudaEventCreate(&start);
cudaEventCreate(&end);
cudaEventRecord(start);

for(row=0; row<m; row++){

    for (d = 0 ; d < m; d++) {

        a[d]=fst[row][d];

        }
        // printf("ok a\n");
        cudaMemcpy(dev_a,a,m*sizeof(int), cudaMemcpyHostToDevice);
    for (col=0; col<m; col++){
        for (d= 0 ; d < m; d++) {

```

```

        b[d]=t_snd[col][d];
        ans[d]=0;

    }
    // printf("ok b\n");
    cudaMemcpy(dev_b,b,m*sizeof(int), cudaMemcpyHostToDevice);
    cudaMemcpy(dev_ans,ans,m*sizeof(int), cudaMemcpyHostToDevice);
    // printf("calling GPU\n");
    mul_r<<<1,m>>>(dev_a,dev_b,dev_ans);
    err=cudaMemcpy(ans,dev_ans,m*sizeof(int), cudaMemcpyDeviceToHost);
    if (err !=cudaSuccess)
    {   printf("failed to copy from device \n");
        exit(EXIT_FAILURE);
    }
    //printf("GPU returned\n");
    //a=fst[0];
    sum_c=0;
    for (d = 0 ; d < m; d++) {

        //printf("%d\t", ans[d]);
        sum_c+=ans[d];
    }

    snd[row][col]=sum_c;

//printf("one element=%d\n",snd[row][col]);
// printf("\n");

}
}
//
cudaEventRecord(end);
cudaEventSynchronize(end);
float time = 0;
cudaEventElapsedTime(&time, start, end);
printf("execution time=%f\n",time);
//
printf(" Matrix multipliation ans=:\n");
for (c = 0; c < n; c++) {
    for (d = 0 ; d < m; d++) {

        printf("%d\t", snd[c][d]);

    }
    printf("\n");
}

//
return 0;
}

```

display the elements of first matrix

| | | | | | | | | | |
|----|---|---|----|----|---|---|----|----|----|
| 4 | 7 | 8 | 6 | 4 | 6 | 7 | 3 | 10 | 2 |
| 3 | 8 | 1 | 10 | 4 | 7 | 1 | 7 | 3 | 7 |
| 2 | 9 | 8 | 10 | 3 | 1 | 3 | 4 | 8 | 6 |
| 10 | 3 | 3 | 9 | 10 | 8 | 4 | 7 | 2 | 3 |
| 10 | 4 | 2 | 10 | 5 | 8 | 9 | 5 | 6 | 1 |
| 4 | 7 | 2 | 1 | 7 | 4 | 3 | 1 | 7 | 2 |
| 6 | 6 | 5 | 8 | 7 | 6 | 7 | 10 | 4 | 8 |
| 5 | 6 | 3 | 6 | 5 | 8 | 5 | 5 | 4 | 1 |
| 8 | 9 | 7 | 9 | 9 | 5 | 4 | 2 | 5 | 10 |
| 3 | 1 | 7 | 9 | 10 | 3 | 7 | 7 | 5 | 10 |

display the elements of second matrix

| | | | | | | | | | |
|----|----|---|----|---|----|---|----|----|----|
| 6 | 1 | 5 | 9 | 8 | 2 | 8 | 3 | 8 | 3 |
| 3 | 7 | 2 | 1 | 7 | 2 | 6 | 10 | 5 | 10 |
| 1 | 10 | 2 | 8 | 8 | 2 | 2 | 6 | 10 | 8 |
| 8 | 7 | 8 | 4 | 7 | 6 | 7 | 4 | 10 | 5 |
| 9 | 2 | 3 | 10 | 4 | 10 | 1 | 9 | 9 | 6 |
| 1 | 10 | 7 | 4 | 9 | 6 | 7 | 2 | 2 | 6 |
| 10 | 9 | 5 | 9 | 2 | 1 | 4 | 1 | 5 | 5 |
| 5 | 5 | 8 | 7 | 4 | 2 | 8 | 6 | 10 | 7 |
| 3 | 2 | 8 | 9 | 6 | 8 | 5 | 2 | 9 | 6 |
| 10 | 8 | 6 | 4 | 9 | 9 | 4 | 2 | 9 | 10 |

Transpose of second Matrix:

| | | | | | | | | | |
|---|----|----|----|----|----|----|----|---|----|
| 6 | 3 | 1 | 8 | 9 | 1 | 10 | 5 | 3 | 10 |
| 1 | 7 | 10 | 7 | 2 | 10 | 9 | 5 | 2 | 8 |
| 5 | 2 | 2 | 8 | 3 | 7 | 5 | 8 | 8 | 6 |
| 9 | 1 | 8 | 4 | 10 | 4 | 9 | 7 | 9 | 4 |
| 8 | 7 | 8 | 7 | 4 | 9 | 2 | 4 | 6 | 9 |
| 2 | 2 | 2 | 6 | 10 | 6 | 1 | 2 | 8 | 9 |
| 8 | 6 | 2 | 7 | 1 | 7 | 4 | 8 | 5 | 4 |
| 3 | 10 | 6 | 4 | 9 | 2 | 1 | 6 | 2 | 2 |
| 8 | 5 | 10 | 10 | 9 | 2 | 5 | 10 | 9 | 9 |
| 3 | 10 | 8 | 5 | 6 | 6 | 5 | 7 | 6 | 10 |

execution time=4.661696

Matrix multipliation ans=:

| | | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 278 | 357 | 303 | 377 | 361 | 261 | 288 | 251 | 428 | 372 |
| 290 | 323 | 301 | 264 | 348 | 268 | 300 | 248 | 389 | 355 |
| 289 | 342 | 287 | 316 | 358 | 263 | 274 | 268 | 451 | 385 |
| 353 | 323 | 330 | 400 | 375 | 295 | 327 | 276 | 456 | 348 |

| | | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 350 | 344 | 352 | 406 | 369 | 266 | 349 | 236 | 439 | 340 |
| 198 | 196 | 186 | 254 | 238 | 205 | 183 | 196 | 274 | 253 |
| 404 | 413 | 374 | 427 | 418 | 319 | 359 | 306 | 526 | 445 |
| 249 | 295 | 265 | 301 | 303 | 218 | 269 | 223 | 341 | 301 |
| 405 | 408 | 342 | 418 | 463 | 360 | 336 | 329 | 535 | 463 |
| 413 | 381 | 345 | 429 | 379 | 345 | 287 | 272 | 525 | 412 |